

- ✓ Name: Brian Ryu

## Computing ID: kfg2ec

```
# Sources used: ChatGPT to help with graphing and formatting with the models created, as well as the function creation

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import ConfusionMatrixDisplay, precision_score, recall_score, accuracy_score
from sklearn.ensemble import RandomForestClassifier

!pip install lime
from lime import lime_tabular
from sklearn.inspection import permutation_importance

# Below is a new package needed for this lab
!pip install ucimlrepo
from ucimlrepo import fetch_ucirepo

Requirement already satisfied: lime in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (0.2.0.1)
Requirement already satisfied: matplotlib in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: numpy in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: scipy in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: tqdm in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: scikit-learn>=0.18 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: scikit-image>=0.12 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: networkx>=3.0 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: pillow>=10.1 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: tifffile>=2022.8.12 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: packaging>=21 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: lazy-loader>=0.4 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: joblib>=1.2.0 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: threadpoolctl>=3.1.0 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: contourpy>=1.0.1 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: cycler>=0.10 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: fonttools>=4.22.0 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: kiwisolver>=1.3.1 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: pyparsing>=3 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: python-dateutil>=2.7 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: six>=1.5 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: ucimlrepo in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (0.0.7)
Requirement already satisfied: pandas>=1.0.0 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: certifi>=2020.12.5 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: numpy>=1.26.0 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: python-dateutil>=2.8.2 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages
Requirement already satisfied: pytz>=2020.1 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: tzdata>=2022.7 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
Requirement already satisfied: six>=1.5 in /Users/brian/.pyenv/versions/3.13.7/lib/python3.13/site-packages (from lime)
```

## ✓ Lab 3: Decision Trees and Random Forests (100 Points)


**Due November 10th, 2025 at 11:59PM**

 Open in Colab

The goal of this lab is to optimize Decision Tree and Random Forest models using the provided dataset on census level data. Your goal is to build a Random Forest Classifier to be able to predict income levels above or below 50k.

The guidance this week is less prescriptive in terms of steps, so use the skills you have gained over the semester to build and evaluate your models. You will be graded on your model building, interpretation of the results and explanation of model selection. As always, you are welcome to rely on your classmates but submit your own code. Lastly, there are likely several correct approaches involving a variety of different conclusions, just make sure your conclusions are supported by your approach.

The dataset should be familiar as it's the census data, on 48,000+ individuals with a variety of variables and a target variable for above or below 50k in salary.



Look through the data dictionary at its source link: <https://archive.ics.uci.edu/ml/datasets/Adult>

## ▼ Part 1: Data Preparation and EDA (15 points)

In a text cell, answer the following exploratory questions and support your observations with any code, if needed.

### ▼ Question 1 (2 points):

Read in the features (X) as a Pandas DataFrame. Show the first 5 rows of the features. How many rows do you have?

```
# Fetch dataset
adult = fetch_uci_repo(id=2)

X = adult.data.features
y = adult.data.targets
```

```
X = pd.DataFrame(X)
display(X.head())
print("Rows:", X.shape[0])
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0

There is a total of... 48,842 rows in features.

## ▼ Question 2 (2 points):

Are there any potential issues in the data or target that need to be corrected? Why are they issues? What specific method would you use to correct them and why?

Consider using code and reading the data description

(<https://archive.ics.uci.edu/dataset/2/adult>) to explore:

- Assumptions and ranges of collected data
- Missing values (impute? drop?)
- Numerical data types represented as strings
- Encoding categorical data appropriately
- Normalization
- Standardization

You will not need to consider feature imbalances or sampling in part 1 or 2 of the lab.

```
X.info()
display(X.head())

for col in X.columns:
    if X[col].dtype == 'object':
        print(f"\nColumn: {col}")
        print(X[col].value_counts().head)
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   age         48842 non-null   int64  
 1   workclass   47879 non-null   object  
 2   fnlwgt     48842 non-null   int64  
 3   education   48842 non-null   object  
 4   education-num 48842 non-null   int64  
 5   marital-status 48842 non-null   object  
 6   occupation   47876 non-null   object  
 7   relationship 48842 non-null   object  
 8   race        48842 non-null   object  
 9   sex         48842 non-null   object  
 10  capital-gain 48842 non-null   int64  
 11  capital-loss 48842 non-null   int64  
 12  hours-per-week 48842 non-null   int64  
 13  native-country 48568 non-null   object  
dtypes: int64(6), object(8)
memory usage: 5.2+ MB
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	40
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40

Column: workclass  
<bound method NDFrame.head of workclass>  
Private 33906  
Self-emp-not-inc 3862  
Local-gov 3136  
State-gov 1981  
? 1836  
Self-emp-inc 1695  
Federal-gov 1432  
Without-pay 21  
Never-worked 10  
Name: count, dtype: int64>

Column: education  
<bound method NDFrame.head of education>  
HS-grad 15784  
Some-college 10878  
Bachelors 8025  
Masters 2657  
Assoc-voc 2061  
11th 1812  
Assoc-acdm 1601  
10th 1389  
7th-8th 955  
Prof-school 834  
9th 756  
12th 657  
Doctorate 594  
5th-6th 509  
1st-4th 247  
Preschool 83  
Name: count, dtype: int64>

Column: marital-status  
<bound method NDFrame.head of marital-status>  
Married-civ-spouse 22379  
Never-married 16117  
Divorced 6633  
Separated 1530  
Widowed 1518  
Married-spouse-absent 628  
Married-AF-spouse 37  
Name: count, dtype: int64>

Column: occupation

```
Column: occupation
<bound method NDFrame.head of occupation>
Prof-specialty      6172
Craft-repair        6112
Exec-managerial     6086
Adm-clerical        5611
Sales                5504
Other-service        4923
Machine-op-inspct   3022
Transport-moving    2355
Handlers-cleaners   2072
?                   1843
Farming-fishing     1490
Tech-support         1446
Protective-serv     983
Priv-house-serv     242
Armed-Forces         15
Name: count, dtype: int64>
```

```
Column: relationship
<bound method NDFrame.head of relationship>
Husband              19716
Not-in-family        12583
Own-child            7581
Unmarried            5125
Wife                 2331
Other-relative       1506
Name: count, dtype: int64>
```

```
Column: race
<bound method NDFrame.head of race>
White                41762
Black                4685
Asian-Pac-Islander   1519
Amer-Indian-Eskimo   470
Other                 406
Name: count, dtype: int64>
```

```
Column: sex
<bound method NDFrame.head of sex>
Male                 32650
Female               16192
Name: count, dtype: int64>
```

```
Column: native-country
<bound method NDFrame.head of native-country>
United-States          43832
Mexico                  951
?                      583
Philippines             295
Germany                 206
Puerto-Rico              184
Canada                  182
El-Salvador              155
India                    151
Cuba                     138
England                  127
China                     122
South                     115
Jamaica                  106
Italy                     105
Dominican-Republic       103
Japan                     92
Guatemala                 88
Poland                     87
Vietnam                     86
Columbia                   85
Haiti                     75
Portugal                   67
Taiwan                     65
Iran                      59
Greece                     49
Nicaragua                   49
Peru                      46
Ecuador                   45
France                     38
Ireland                   37
Hong                      30
Thailand                   30
Cambodia                   28
Trinidad&Tobago           27
Laos                      23
Yugoslavia                 23
Outlying-US(Guam-USVI-etc) 23
Scotland                   21
Name: count, dtype: int64>
```

Honuras

**Potential issues:** Potential issues I've identified are.... missing values that have been substituted by "?". When looking into columns such as `workclass`, there is a <sup>20</sup> `count` with 1836, which makes it difficult to see what this value is representing (does not use NaN). The way I would handle this would be to put NaN instead of "?" and then dropping the missing rows if the values are minimal such as less than 5%. Another issue could be to one-hot encode categorical variables because tree models require this type of action. Another potential issue to put down could be to look for possible outliers such as having an age of 1000 which is not realistic at all. For y, we can also clean this by making sure that there are no trailing periods for target labels such as '>50K'.

### ▼ Question 3 (6 points):

Preprocess the data according to the issues and correction methods you've identified. Save the new features and target variable (if necessary) as X\_clean and y\_clean.

```
X = X.replace("?", np.nan)

X_clean = X.dropna()

X_clean = pd.get_dummies(X_clean, drop_first=True)

y_clean = y.apply(lambda s: s.str.strip().replace('.', ''))

display(y_clean.head())

display(X_clean.head())
```

#### income

0	<=50K
1	<=50K
2	<=50K
3	<=50K
4	<=50K

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	workclass_Local-gov	workclass_Private	workclass_Self-emp-inc	workclass_emp-no
0	39	77516	13	2174	0	40	False	False	False	False
1	50	83311	13	0	0	13	False	False	False	False
2	38	215646	9	0	0	40	False	True	False	False
3	53	234721	7	0	0	40	False	True	False	False
4	28	338409	13	0	0	40	False	True	False	False

5 rows x 96 columns

### ▼ Question 4 (5 points):

Create 2 versions of y\_clean to create a new target response of whether income is above or below \$50,000 for classification.

1. **y\_clean\_binary:** Recode y\_clean to be 1 if the target is over or equal to 50,000
2. **y\_clean\_string:** Recode y\_clean to be "Above or Equal" if the target is over or equal to 50,000 and "Below" if under.

```
if isinstance(y_clean, pd.DataFrame):
    y_clean = y_clean.iloc[:, 0]

y_clean = y_clean[X_clean.index]

y_clean_binary = y_clean.apply(lambda s: 1 if s == '>50K' else 0)
y_clean_string = y_clean.apply(lambda s: "Above or Equal" if s == '>50K' else "Below")
```

```

print(y_clean_binary.value_counts())
print(y_clean_string.value_counts())

income
0    37714
1     7508
Name: count, dtype: int64
income
Below        37714
Above or Equal    7508
Name: count, dtype: int64

```

## ▼ Part 2: Decision Tree Pruning, Tuning and Evaluation (30 Points)

### ▼ Question 1 (5 points):

Create a function to take in a feature variable (X) and (y). In this function, do the following:

- Create a train test split with a random seed of 3001.
- Use a vanilla decision tree model to fit the model on the train set and predict on the test set.
- Print the precision, recall, and accuracy of the model after prediction.

Test your function on both y\_clean\_binary and y\_clean\_string. For any of the following questions, you may use whichever y\_clean variable you'd like.

```

def eval_dt(X, y, max_depth=None):
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.25, random_state=3001)

    dt = DecisionTreeClassifier(random_state=3001, max_depth=max_depth)
    dt.fit(X_train, y_train)

    y_pred = dt.predict(X_test)

    if y.dtype.kind in ['O', 'U', 'S']:
        pos_label = "Above or Equal"
    else:
        pos_label = 1

    precision = precision_score(y_test, y_pred, pos_label=pos_label)
    recall = recall_score(y_test, y_pred, pos_label=pos_label)
    accuracy = accuracy_score(y_test, y_pred)

    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"Accuracy: {accuracy:.4f}")

    plt.figure(figsize=(20,10))
    plot_tree(dt, feature_names=X.columns, class_names=[str(c) for c in dt.classes_],
              filled=True, fontsize=8, max_depth=5)
    plt.title("Decision Tree Visualization")
    plt.show()

    cm_display = ConfusionMatrixDisplay.from_estimator(dt, X_test, y_test, display_labels=[str(c) for c in dt.classes_])
    cm_display.ax_.set_title("Confusion Matrix")
    plt.show()

    return dt

print("Using y_clean_binary:")
eval_dt(X_clean, y_clean_binary, max_depth=5)



print("\nUsing y_clean_string:")
eval_dt(X_clean, y_clean_string, max_depth=5)

dt = eval_dt(X_clean, y_clean_binary, max_depth=5)


```




Using `y_clean_binary`:  
 Precision: 0.6541  
 Recall: 0.2554  
 Accuracy: 0.8519




Using `y_clean_string`:  
 Precision: 0.6541  
 Recall: 0.2554  
 Accuracy: 0.8519





Confusion Matrix




Precision: 0.6541


Recall: 0.2554

Accuracy: 0.8519

Decision Tree Visualization



Confusion Matrix



**Question 2 (5 points):**0 1  
1000

Adjust your function to include some plotting features. After your prediction code, plot:

1. A visualization of the resulting Decision Tree
2. A confusion matrix of the results

Your tree might be overwhelming or very large! If it is too large to be interpreted, constrain the max\_depth parameter manually to 5 or less.

Graphs of the resulting decision trees and confusion matrices are displayed above.

**Question 3 (5 points):**

Create a sorted list of feature importances and comment on the top features. Plot your feature importances in a horizontal or vertical bar chart from most to least important. Label each bar with its feature importance rounded to the nearest integer (ie: 30%).

Are there a few that seem to be more important than the others?

```
feature_importances = dt.feature_importances_
features = X_clean.columns

importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)


importance_df = importance_df.head(20)
importance_df['Importance (%)'] = (importance_df['Importance'] * 100).round(0)

print("Sorted Feature Importance")
print(importance_df)

plt.figure(figsize=(10, 6))
plt.barh(importance_df['Feature'], importance_df['Importance (%)'], color='skyblue')
plt.gca().invert_yaxis()
plt.xlabel('Importance (%)')
plt.title('Decision Tree Feature Importances')

for i, v in enumerate(importance_df['Importance (%):']):
    plt.text(v + 0.5, i, f"{int(v)}%", va='center')

plt.show()
```



Yes, the ones that seem to be more important than the others consist of a few such as the marital-status\_Married-civ-spouse, education-num, capital-gain as they all have above 20%. The rest of the data are either less than 5% or just 0%.

#### ▼ Question 4 (5 points):

Write at least 5 sentences interpreting the results of your decision tree, confusion matrix, and feature importance visualizations.

Is there any aspect of your results that you are uncertain or unsure of?

**My results show...** that the decision tree achieved a pretty good level of accuracy in predicting the individual's income of if it is above or below \$50,000. From the confusion matrix, the model does seem to perform well overall but sometimes misclassify some individuals with lower incomes as higher, possibly showing an imbalance or overlap in features between two classes. The top features that were possibly identified was education level, capital gain, and hours worked per week, which makes realistic sense. In general, individuals with higher education, more capital gains, and longer working hours lead to more money being made to higher income.

For an aspect of my result that I am uncertain or unsure of is the fact that the results are from the one-hot encoded categorical variables. This is because doing so may lead to some being more dominant in importance just because they had more representation or unique

splits. There is also a vulnerability of decision trees being overfitted when it is not pruned or `max_depth` is set high. This makes it so that the deeper the trees are, the more unreliable it can become. The values from the precision and recall may hint at something like a Random Forest being better in generalizing.

## ▼ Question 5 (5 points):

Finally, we will create a new function to tune your decision tree to get more accurate and efficient results. Update your function to take in several new parameters with these default values:

- `criterion_val = 'gini'`
- `splitter_val = 'best'`
- `max_depth_val = None`
- `min_samples_split_val = 2`
- `min_samples_leaf_val = 1`

Pass your own variable into the decision tree by specifying what `sklearn` parameter you are trying to tune. This will simply be the parameter without the "`_val`" suffix.

**For example, if your vanilla decision tree variable is called `clf`, you would adjust it like this:**

```
clf = DecisionTreeClassifier(criterion=criterion_val, splitter=splitter_val, ...)
```

```
def eval_dtreet_tuned(X, y,
                      criterion_val='gini',
                      splitter_val='best',
                      max_depth_val=None,
                      min_samples_split_val=2,
                      min_samples_leaf_val=1
                      ):
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.25, random_state=3001
    )

    clf = DecisionTreeClassifier(
        criterion=criterion_val,
        splitter=splitter_val,
        max_depth=max_depth_val,
        min_samples_split=min_samples_split_val,
        min_samples_leaf=min_samples_leaf_val,
        random_state=3001
    )

    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)

    if y.dtype == 'object':
        pos_label = "Above or Equal"
    else:
        pos_label = 1

    precision = precision_score(y_test, y_pred, pos_label=pos_label, zero_division=0)
    recall = recall_score(y_test, y_pred, pos_label=pos_label, zero_division=0)
    accuracy = accuracy_score(y_test, y_pred)

    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"Accuracy: {accuracy:.4f}")

    plt.figure(figsize=(18, 8))
    plot_tree(clf, filled=True, max_depth=3, fontsize=6)
    plt.title("Decision Tree Visualization (max_depth=3 for readability)")
    plt.show()

    ConfusionMatrixDisplay.from_predictions(y_test, y_pred, cmap='Blues')
    plt.title("Confusion Matrix")
```

```
plt.show()  
return clf
```

## ▼ Question 6 (5 points):

Call your new function with either clean y variable at least 3 times. Each time, vary the values for all the parameters and examine its effects on your tree, confusion matrix, and metrics.

You will likely want to look at documentation to see accepted values: <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

Why did you pick the values you did? What combination had the best effect on accuracy? Were you surprised by any of the results?

```
print("Model 1: Default")  
dt1 = eval_dtreetuned(  
    X_clean, y_clean_binary,  
    criterion_val='gini',  
    splitter_val='best',  
    max_depth_val=None,  
    min_samples_split_val=2,  
    min_samples_leaf_val=1  
)  
  
print("Model 2: changed to entropy, moderate pruning")  
dt2 = eval_dtreetuned(  
    X_clean, y_clean_binary,  
    criterion_val='entropy',  
    splitter_val='best',  
    max_depth_val=5,  
    min_samples_split_val=10,  
    min_samples_leaf_val=4  
)  
  
print("\nModel 3: random split, more and more and more pruning")  
dt3 = eval_dtreetuned(  
    X_clean, y_clean_binary,  
    criterion_val='gini',  
    splitter_val='random',  
    max_depth_val=3,  
    min_samples_split_val=20,  
    min_samples_leaf_val=8  
)
```